

# **Week 1: Data Collection for ML**

**CS 203: Software Tools and Techniques for AI**

Prof. Nipun Batra, IIT Gandhinagar

# Your Mission



You work at Netflix.

Your boss asks:

"Which movies should we add to our catalog next month?"

This is an ML problem.

But first... we need data.

# The ML Pipeline

```
DATA → MODEL → PREDICTIONS  
↑  
We are here
```

**Today's question:** Where does data come from?

# What Data Do We Need?

To predict movie success, we need features:

-  **Title, Genre, Year**
-  **IMDb Rating**
-  **Rotten Tomatoes Score**
-  **Box Office Performance**
-  **Cast & Director**
-  **Social Media Buzz**

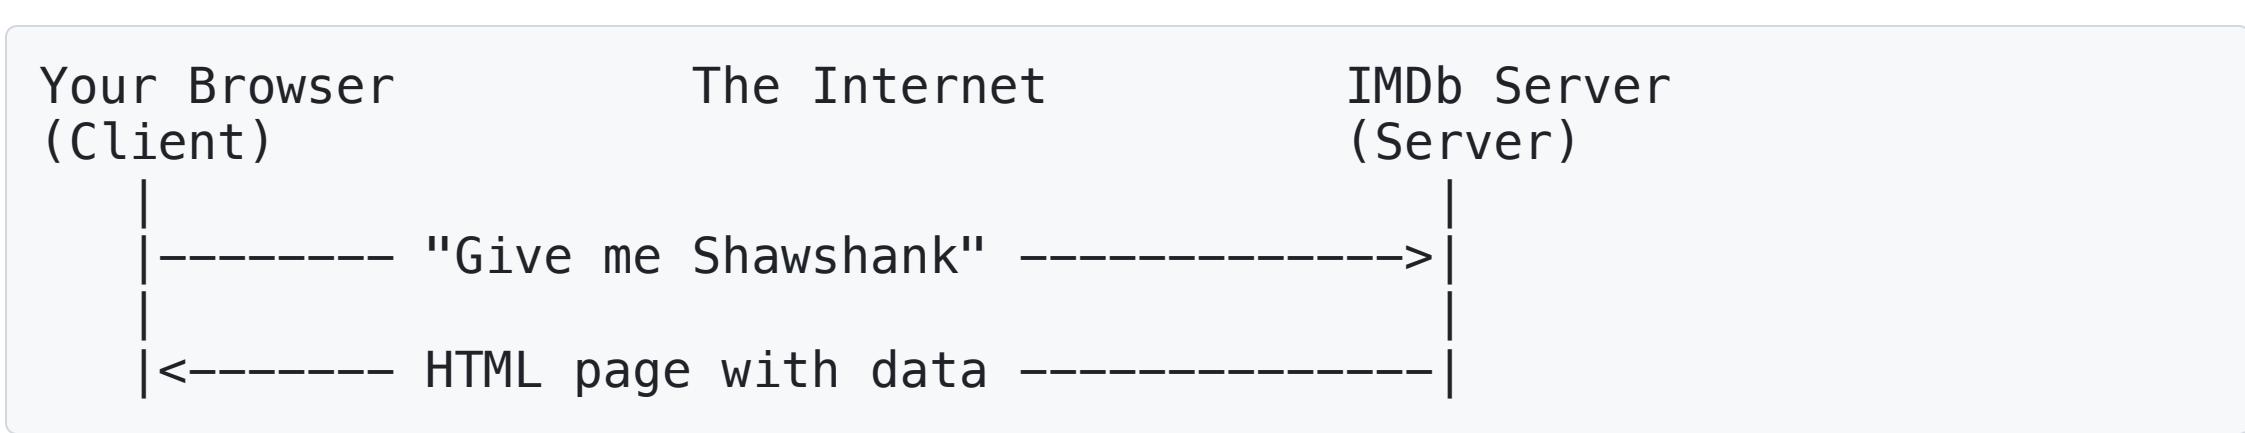
**Problem:** This data lives on websites!

# Today's Agenda

1. **Understanding the Web** (HTTP, URLs, requests)
2. **Tool #1: curl** (command-line HTTP)
3. **Tool #2: Chrome DevTools** (inspect websites)
4. **Tool #3: Python requests** (automate data fetching)
5. **Tool #4: BeautifulSoup** (parse HTML)
6. **Tool #5: Playwright** (dynamic websites)
7. **APIs: The Better Way** (structured data)

# **Part 1: How the Web Works**

# Client-Server Model



**HTTP** = HyperText Transfer Protocol

The language browsers and servers speak.

# A Simple HTTP Request

URL: `https://www.imdb.com/title/tt0111161/`

Breaking it down:

- `https://` → Protocol (secure HTTP)
- `www.imdb.com` → Domain (server)
- `/title/tt0111161/` → Path (resource)

# HTTP Methods

**GET** → Retrieve data

```
GET /title/tt0111161/
```

**POST** → Send data

```
POST /api/login  
Body: {"username": "user", "password": "pass"}
```

**Others:** PUT, DELETE, PATCH

*Today we focus on GET*

# HTTP Response

Server sends back:

**Status Code:**

- 200 Success
- 404 Not found
- 500 Server error
- 429 Rate limited

**Headers:** Metadata (content type, length)

**Body:** The actual data (HTML, JSON, etc.)

## **Part 2: Tool #1 - curl**

# What is curl?

**Command-line tool** to make HTTP requests.

Pre-installed on Mac/Linux.

Windows: Use Git Bash or WSL.

# curl Basics

Get IMDb homepage:

```
curl https://www.imdb.com/
```

Output: Raw HTML (lots of it!)

# curl with Options

**Save to file:**

```
curl https://www.imdb.com/ -o imdb.html
```

**See response headers:**

```
curl -I https://www.imdb.com/
```

**Follow redirects:**

```
curl -L https://imdb.com/
```

# curl for APIs

Get movie data (JSON):

```
curl "http://www.omdbapi.com/?apikey=YOUR_KEY&t=Shawshank"
```

Pretty print JSON:

```
curl "http://www.omdbapi.com/?apikey=YOUR_KEY&t=Shawshank" | jq
```

*jq = JSON processor*

# Why Learn curl?

-  Quick testing of APIs
-  No code needed
-  Works everywhere (servers, containers)
-  Great for debugging

## **Part 3: Tool #2 - Chrome DevTools**

# Open DevTools

1. Visit: <https://www.imdb.com/title/tt0111161/>
2. Right-click → "Inspect"
3. Go to **Network** tab
4. Refresh page (Cmd+R / Ctrl+R)

# What You See

**Network Tab** shows:

- All HTTP requests the page makes
- Status codes
- Response times
- Headers
- Response data

Demo time! 

# Inspecting Elements

## Elements Tab:

- Shows HTML structure
- Click element → highlights on page
- Right-click element on page → Inspect

## Finding data:

- Look for rating: `<span class="rating">9.3</span>`
- Note the class name: `rating`
- We'll use this later!

# DevTools for APIs

Many websites load data via **background API calls**.

**Network tab → Fetch/XHR filter**

Shows JSON responses!

**Example:** Twitter, Reddit, YouTube

## **Part 4: Tool #3 - Python requests**

# Why Python?

curl is great for testing.

But for **automation**, we need code.

Python `requests` library = curl for Python

# Install requests

```
pip install requests
```

That's it! 

# Basic GET Request

```
import requests

response = requests.get('https://www.imdb.com/title/tt0111161/')

print(response.status_code) # 200
print(response.text[:100]) # First 100 chars of HTML
```

# Response Object

```
response = requests.get('https://www.imdb.com/')

# Status
response.status_code # 200

# Headers
response.headers['content-type'] # 'text/html'

# Body
response.text # HTML as string
response.content # HTML as bytes
```

# Query Parameters

Instead of:

```
url = 'https://api.example.com/search?q=movie&year=2024'
```

Better:

```
response = requests.get('https://api.example.com/search',
    params={
        'q': 'movie',
        'year': 2024
    }
)
```

# Headers

Send custom headers:

```
headers = {  
    'User-Agent': 'Mozilla/5.0 (Netflix Bot)',  
    'Accept': 'application/json'  
}  
  
response = requests.get(url, headers=headers)
```

Why? Some sites block Python's default User-Agent.

# Error Handling

```
response = requests.get(url)

if response.status_code == 200:
    print("Success!")
    data = response.text
elif response.status_code == 404:
    print("Page not found")
else:
    print(f"Error: {response.status_code}")
```

# Better Error Handling

```
try:  
    response = requests.get(url, timeout=10)  
    response.raise_for_status() # Raises exception for 4xx/5xx  
    data = response.text  
except requests.exceptions.Timeout:  
    print("Request timed out")  
except requests.exceptions.HTTPError as e:  
    print(f"HTTP error: {e}")  
except requests.exceptions.RequestException as e:  
    print(f"Error: {e}")
```

## **Part 5: Tool #4 - BeautifulSoup**

# The Problem

We got the HTML:

```
<html>
  <body>
    <h1 class="title">The Shawshank Redemption</h1>
    <span class="rating">9.3</span>
    <div class="genre">Drama</div>
  </body>
</html>
```

How do we **extract** the rating?

# BeautifulSoup

HTML parser for Python.

Converts messy HTML → searchable structure.

```
pip install beautifulsoup4
```

# Basic Usage

```
from bs4 import BeautifulSoup
import requests

# Get the page
response = requests.get('https://www.imdb.com/title/tt0111161/')

# Parse HTML
soup = BeautifulSoup(response.text, 'html.parser')

# Now we can search!
```

# Finding Elements

By tag:

```
title = soup.find('h1')
print(title.text) # "The Shawshank Redemption"
```

By class:

```
rating = soup.find('span', class_='rating')
print(rating.text) # "9.3"
```

By ID:

```
element = soup.find(id='main-content')
```

# Find vs Find\_all

`find()` → First match

```
first_div = soup.find('div')
```

`find_all()` → All matches (returns list)

```
all_divs = soup.find_all('div')
print(len(all_divs)) # e.g., 42

for div in all_divs:
    print(div.text)
```

# CSS Selectors

More powerful searching:

```
# Find all links in a specific div  
links = soup.select('div.cast-list a')  
  
# Complex selector  
rating = soup.select_one('div.rating-container span.value')
```

**Tip:** Copy selector from Chrome DevTools!

# Extracting Attributes

```
# Get link URL
link = soup.find('a', class_='movie-link')
url = link['href']
print(url) # "/title/tt0111161/"

# Get image source
img = soup.find('img', class_='poster')
poster_url = img['src']
```

# Netflix Example: Complete Scraper

```
import requests
from bs4 import BeautifulSoup

def get_movie_info(imdb_id):
    url = f'https://www.imdb.com/title/{imdb_id}/'
    response = requests.get(url)
    soup = BeautifulSoup(response.text, 'html.parser')

    title = soup.find('h1').text.strip()
    rating = soup.find('span', class_='rating').text
    genre = soup.find('div', class_='genre').text.strip()

    return {
        'title': title,
        'rating': float(rating),
        'genre': genre
    }

movie = get_movie_info('tt0111161')
print(movie)
```

# Web Scraping Challenges

- ✗ HTML structure changes → code breaks
- ✗ Slow (one request at a time)
- ✗ Anti-scraping measures (CAPTCHAs, blocks)
- ✗ Legal/ethical concerns
- ✗ Dynamic content (JavaScript-loaded)

## **Part 6: Tool #5 - Playwright**

# The JavaScript Problem

Modern websites load data with JavaScript.

```
<!-- Initial HTML (empty!) -->
<div id="movies"></div>

<script>
  // JavaScript loads movies after page loads
  fetch('/api/movies').then(data => {
    document.getElementById('movies').innerHTML = data;
  });
</script>
```

requests + BeautifulSoup only sees empty <div> !

# Solution: Headless Browser

**Playwright** = Control a real browser from Python

- Executes JavaScript
- Waits for content to load
- Can click, scroll, type

# Install Playwright

```
pip install playwright  
playwright install chromium
```

Downloads a Chromium browser.

# Basic Playwright Usage

```
from playwright.sync_api import sync_playwright

with sync_playwright() as p:
    # Launch browser
    browser = p.chromium.launch(headless=True)
    page = browser.new_page()

    # Navigate to page
    page.goto('https://www.imdb.com/title/tt0111161/')

    # Wait for content
    page.wait_for_selector('.rating')

    # Get HTML (after JavaScript runs)
    html = page.content()

    browser.close()
```

# Playwright Features

**Click elements:**

```
page.click('button.load-more')
```

**Fill forms:**

```
page.fill('input[name="search"]', 'Shawshank')
page.press('input[name="search"]', 'Enter')
```

**Screenshot:**

```
page.screenshot(path='screenshot.png')
```

# When to Use Playwright?

✓ JavaScript-heavy sites (React, Angular, Vue)

✓ Need to interact (click, scroll, login)

✓ Content loads on scroll (infinite scroll)

✗ Slow and heavy (uses real browser)

✗ Overkill for static sites

**Rule:** Try requests first. Use Playwright if needed.

# Netflix Example: Dynamic Scraping

```
from playwright.sync_api import sync_playwright
from bs4 import BeautifulSoup

def scrape_dynamic_site(url):
    with sync_playwright() as p:
        browser = p.chromium.launch()
        page = browser.new_page()
        page.goto(url)

        # Wait for JavaScript to load content
        page.wait_for_selector('.movie-list')

        # Get the fully-loaded HTML
        html = page.content()
        browser.close()

        # Parse with BeautifulSoup
        soup = BeautifulSoup(html, 'html.parser')
        movies = soup.find_all('div', class_='movie-card')
    return [m.text for m in movies]
```

## **Part 7: APIs - The Better Way**

# The Problem with Scraping

We've learned scraping but...

- 😞 HTML changes break code
- 🐌 Slow
- 🚫 Sites may block you
- ⚖️ Legal gray area

**Better solution:** Use an API!

# What is an API?

**API** = Application Programming Interface

**Web API** = URL that returns structured data (usually JSON)

Instead of:

HTML (for humans) → scrape → extract data

API gives:

JSON (for machines) → directly usable data

# HTML vs JSON

HTML (for browsers):

```
<div class="movie">
  <h1>Shawshank Redemption</h1>
  <span class="rating">9.3</span>
</div>
```

JSON (for code):

```
{
  "title": "Shawshank Redemption",
  "rating": 9.3
}
```

Much cleaner! ✨

# REST APIs

**REST** = Representational State Transfer

Common pattern for web APIs.

**Example endpoints:**

GET /movies	→ List all movies
GET /movies/123	→ Get movie #123
POST /movies	→ Create new movie
PUT /movies/123	→ Update movie #123
DELETE /movies/123	→ Delete movie #123

# **OMDb API**

**The Open Movie Database**

Free API for movie data!

Website: <http://www.omdbapi.com/>

**Get API key:** Sign up (free)

# Using OMDb API

```
import requests

API_KEY = 'your_key_here'

response = requests.get('http://www.omdbapi.com/', params={
    'apikey': API_KEY,
    't': 'Shawshank Redemption' # Search by title
})

data = response.json() # Parse JSON → Python dict
print(data['Title']) # "The Shawshank Redemption"
print(data['imdbRating']) # "9.3"
print(data['Genre']) # "Drama"
```

# JSON Response

```
{  
  "Title": "The Shawshank Redemption",  
  "Year": "1994",  
  "Rated": "R",  
  "Genre": "Drama",  
  "Director": "Frank Darabont",  
  "Actors": "Tim Robbins, Morgan Freeman",  
  "imdbRating": "9.3",  
  "imdbID": "tt0111161",  
  "BoxOffice": "$28,767,189"  
}
```

Clean, structured, perfect for ML! 

# Search by IMDb ID

```
response = requests.get('http://www.omdbapi.com/', params={  
    'apikey': API_KEY,  
    'i': 'tt0111161' # IMDb ID  
})  
  
data = response.json()
```

More reliable than searching by title.

# API Authentication

Most APIs require **authentication**.

**Common methods:**

1. **API Key** (in URL or header)

```
params={'apikey': 'abc123'}
```

2. **Bearer Token** (in header)

```
headers={'Authorization': 'Bearer xyz789'}
```

3. **OAuth** (complex, for user data)

# Rate Limiting

APIs limit requests to prevent abuse.

**Example:** 1000 requests/day

**HTTP 429:** Too Many Requests

**Solutions:**

- Respect limits
- Cache responses
- Pay for higher tier

# Handling Rate Limits

```
import time

def get_movie_safe(movie_id):
    response = requests.get(url, params={...})

    if response.status_code == 429:
        print("Rate limited! Waiting...")
        time.sleep(60) # Wait 1 minute
        return get_movie_safe(movie_id) # Retry

    return response.json()
```

*Better: Use backoff library*

# Pagination

APIs often return data in **pages**.

```
page = 1
all_movies = []

while True:
    response = requests.get(url, params={
        'page': page,
        'limit': 100
    })

    movies = response.json()['results']
    if not movies:
        break # No more data

    all_movies.extend(movies)
    page += 1
```

# Netflix Complete Example

```
import requests

API_KEY = 'your_omdb_key'

def fetch_movie_data(titles):
    """Fetch data for multiple movies."""
    movies = []

    for title in titles:
        response = requests.get('http://www.omdbapi.com/',
            params={'apikey': API_KEY, 't': title})
        )

        if response.status_code == 200:
            data = response.json()
            if data['Response'] == 'True':
                movies.append({
                    'title': data['Title'],
                    'year': data['Year'],
                    'rating': data['imdbRating'],
                    'genre': data['Genre'],
                    'box_office': data.get('BoxOffice', 'N/A')
                })
            else:
                print(f"Movie '{title}' not found or failed to fetch data.")

    return movies
```

# Other Useful APIs

## Movies:

- TMDb (The Movie Database)
- Rotten Tomatoes

## Weather:

- OpenWeatherMap
- Weather.gov

## News:

- NewsAPI
- Guardian API

## Finance:

# API Best Practices

-  **Read the docs** (rate limits, auth)
-  **Handle errors** (network, rate limits, invalid data)
-  **Cache responses** (don't re-fetch same data)
-  **Respect rate limits** (be a good citizen)
-  **Keep API keys secret** (use environment variables)

# Storing API Keys Safely

✗ DON'T:

```
API_KEY = 'abc123xyz' # Hardcoded!
```

✓ DO:

```
import os
API_KEY = os.environ['OMDB_API_KEY']
```

Set in terminal:

```
export OMDB_API_KEY='abc123xyz'
```

Or use `.env` file + `python-dotenv`

# Summary: Tools We Learned

Tool	Purpose	When to Use
<code>curl</code>	Test HTTP from terminal	Quick API testing
<code>Chrome DevTools</code>	Inspect web traffic	Find data sources
<code>requests</code>	HTTP in Python	Static sites, APIs
<code>BeautifulSoup</code>	Parse HTML	Web scraping
<code>Playwright</code>	Control browser	JavaScript sites

# Web Scraping vs APIs

Aspect	Web Scraping	APIs
Speed	Slow	Fast ⚡
Reliability	Fragile 🐛	Stable ✅
Data Format	HTML (messy)	JSON (clean)
Legal	Gray area ⚖️	Approved ✅
When to use	No API available	Always prefer!

**Golden Rule:** Use API if available. Scrape as last resort.

# The Data Formats

**HTML:** For humans (browsers)

```
<div class="price">$29.99</div>
```

**JSON:** For machines (APIs)

```
{"price": 29.99, "currency": "USD"}
```

**CSV:** Tabular data

```
title, rating, year  
Shawshank, 9.3, 1994
```

# Netflix Project Status

 We can now collect movie data!

## Next steps:

1. **Week 2:** Validate the data (Pydantic)
2. **Week 3:** Enrich with LLM features
3. **Week 7:** Build prediction model
4. **Week 9:** Deploy as interactive demo

# Lab Preview

**Your task:** Build a movie dataset collector

1. Use OMDb API to fetch 100 movies
2. Parse and structure the data
3. Save to CSV
4. Handle errors gracefully

**Bonus:** Compare with web scraping approach

# Key Takeaways

- 🎯 Data collection is the first step in ML
- 🔧 Master the tools: curl, DevTools, requests, BeautifulSoup, Playwright
- 📊 APIs > Scraping (when available)
- ⚡ Handle errors (network, rate limits, bad data)
- 🔒 Keep credentials safe (environment variables)

# Resources

## Documentation:

- requests: <https://requests.readthedocs.io/>
- BeautifulSoup: <https://www.crummy.com/software/BeautifulSoup/>
- Playwright: <https://playwright.dev/python/>

## Practice APIs:

- OMDB: <http://www.omdbapi.com/>
- JSONPlaceholder: <https://jsonplaceholder.typicode.com/>
- PokéAPI: <https://pokeapi.co/>

# Questions?

**Next Week: Data Validation with Pydantic**

Making sure our Netflix data is clean and usable!