Week 5: Data Augmentation

CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra

IIT Gandhinagar

Part 1: The Data Hunger Problem

More data from existing data

Previously on CS 203...

Week 1: Collected 10,000 movie records from OMDB API

Week 2: Validated and cleaned the data

Week 3: Labeled 5,000 movies as "good" or "bad"

Week 4: Optimized labeling with active learning + weak supervision

```python
# Current state of our Netflix movie project
labeled_movies = 5000
model_accuracy = 0.82  # 82% accuracy

# But Netflix wants 90%+ accuracy!
# And we've exhausted our labeling budget...
```

Can we improve without more labeling?

The Data Hunger Problem

Deep learning models need data:

- ResNet-50 trained on 1.2M ImageNet images

- GPT-3 trained on 45TB of text

- AlphaGo trained on 30M game positions

Your reality:

- 500 labeled images

- 1,000 text samples

- 100 audio clips

Solution: Create more data from existing data through augmentation

What is Data Augmentation?

Data Augmentation: Apply transformations to existing data to create new training examples

Key Idea: Generate variations that preserve the label but increase diversity

Example (Image):

- Original: Cat image

- Rotated 10°: Still a cat

- Flipped horizontally: Still a cat

- Slightly darker: Still a cat

Benefits:

- More training data without labeling

- Better generalization

- Reduced overfitting

The Photographer Analogy

Imagine you only have ONE photo of a cat to teach someone "what is a cat":

```
One photo: Person might think "cat" means
- This specific pose
- This specific lighting
- This specific background
- This specific angle

Many photos: Person learns
- Cats can be in different poses
- Cats look similar in different lighting
- Cats can be anywhere
- Cats look similar from different angles
```

Augmentation = Taking many "virtual photos" from one real photo!

Free Data: The Augmentation Magic

```python
# Before augmentation
original_dataset = 1000  # Labeled examples
training_epochs = 100

# Each epoch: model sees 1000 examples (same ones!)
# Model memorizes specific examples = OVERFITTING

# After augmentation
augmented_dataset = 1000 * 10  # 10 variations each
training_epochs = 100

# Each epoch: model sees 10,000 DIFFERENT examples!
# Model learns general patterns = GENERALIZATION
```

10x more data for FREE (no labeling cost)!

Why Data Augmentation Works

1. Implicit Regularization

- Model sees slightly different versions

- Learns robust features

- Reduces overfitting

2. Invariance Learning

- Model learns that rotations don't change identity

- Small color shifts don't matter

- Position in frame doesn't change class

3. Coverage of Data Distribution

- Fills gaps in training data

- Simulates real-world variations

- Tests robustness

Data Augmentation vs Data Collection

Data Collection:

- Time: Weeks to months

- Cost: High (labeling, storage)

- Effort: Manual collection and annotation

- Diversity: Limited by budget

Data Augmentation:

- Time: Minutes to hours

- Cost: Low (just compute)

- Effort: Automated transformations

- Diversity: Programmatically generated

Best Practice: Do both! Augmentation complements collection.

Why Image Augmentation Works So Well

Key insight: Geometric changes don't change what's in the image!

```
Original:          Flipped:          Rotated:

┌─────────┐       ┌─────────┐       ┌─────────┐
│  /\_/\  │       │  /\_/\  │       │  /\_/\  │
│ ( o.o ) │  →    │ ( o.o ) │  →    │ ( o.o ) │
│  > ^ <  │       │  > ^ <  │       │  > ^ <  │
└─────────┘       └─────────┘       └─────────┘

It's still a cat! The label doesn't change.
```

This is called "invariance" - the label is invariant to these transforms.

Image Augmentation: Geometric Transforms

Basic transformations:

1. **Rotation**: Rotate ±15-30 degrees

2. **Horizontal Flip**: Mirror image left-right

3. **Vertical Flip**: Mirror image top-bottom (use carefully)

4. **Translation**: Shift image by pixels

5. **Scaling**: Zoom in/out

6. **Shearing**: Skew image

7. **Cropping**: Random crops

Implementation with PIL:

```python
from PIL import Image

img = Image.open('cat.jpg')
rotated = img.rotate(15)
flipped = img.transpose(Image.FLIP_LEFT_RIGHT)
```

The Movie Poster Example

```python
# For our Netflix movie poster classification

from PIL import Image
import albumentations as A

# Original movie poster (e.g., for "Inception")
poster = Image.open("inception_poster.jpg")
label = "Sci-Fi/Thriller"

# Augmented versions
transform = A.Compose([
    A.HorizontalFlip(p=0.5),            # Poster still shows same movie
    A.RandomBrightnessContrast(p=0.3), # Like different lighting
    A.Rotate(limit=10),                 # Slight tilt
])

# Generate 10 variations
for i in range(10):
    augmented = transform(image=np.array(poster))['image']
    # All 10 are still "Inception" posters!
```

Same poster, 10 training examples!

Image Augmentation: Color Transforms

Color space adjustments:

1. **Brightness**: Make lighter/darker

2. **Contrast**: Increase/decrease contrast

3. **Saturation**: Make more/less colorful

4. **Hue**: Shift color spectrum

5. **Grayscale**: Convert to black and white

6. **Color Jittering**: Random color variations

```python
from PIL import ImageEnhance

enhancer = ImageEnhance.Brightness(img)
brighter = enhancer.enhance(1.5)  # 50% brighter


enhancer = ImageEnhance.Contrast(img)
higher_contrast = enhancer.enhance(1.3)
```

Image Augmentation: Advanced Techniques

1. Cutout: Remove random patches

```python
# Remove 16x16 patch
x, y = random.randint(0, w-16), random.randint(0, h-16)
img[y:y+16, x:x+16] = 0
```

2. Mixup: Blend two images

```python
lambda_val = np.random.beta(alpha, alpha)
mixed = lambda_val * img1 + (1 - lambda_val) * img2
label = lambda_val * label1 + (1 - lambda_val) * label2
```

3. CutMix: Replace patch with another image
4. AugMix: Apply multiple augmentations and mix

Albumentations Library

Fast and flexible image augmentation library

```python
import albumentations as A
from albumentations.pytorch import ToTensorV2

transform = A.Compose([
    A.RandomRotate90(),
    A.Flip(),
    A.Transpose(),
    A.GaussNoise(),
    A.OneOf([
        A.MotionBlur(p=0.2),
        A.MedianBlur(blur_limit=3, p=0.1),
        A.Blur(blur_limit=3, p=0.1),
    ], p=0.2),
    A.ShiftScaleRotate(shift_limit=0.0625, scale_limit=0.2, rotate_limit=45, p=0.2),
    A.OneOf([
        A.OpticalDistortion(p=0.3),
        A.GridDistortion(p=0.1),
    ], p=0.2),
    A.Normalize(mean=(0.485, 0.456, 0.406), std=(0.229, 0.224, 0.225)),
    ToTensorV2(),
])
```

Albumentations - Key Features

Why Albumentations?

1. **Fast**: Optimized with NumPy/OpenCV

2. **Flexible**: Easy to compose transformations

3. **Framework-agnostic**: Works with PyTorch, TensorFlow, etc.

4. **Preserves Bounding Boxes**: For object detection

5. **Keypoint Support**: For pose estimation

Common Augmentations:

- Geometric: Rotate, Flip, Shift, Scale

- Blur: Motion, Gaussian, Median

- Noise: Gaussian, ISO, Salt & Pepper

- Weather: Rain, Fog, Snow, Sun Flare

- Advanced: Cutout, CoarseDropout

Image Augmentation Best Practices

1. Choose Appropriate Augmentations

- Natural images: Rotation, flip, color jitter

- Medical images: Be careful with flips (anatomy matters)

- Text/OCR: No rotation, no flip (orientation matters)

2. Augmentation Strength

- Start mild, increase gradually

- Too strong: Model learns wrong patterns

- Too weak: No benefit

3. Validation Set

- Don't augment validation/test data

- Measure performance on real distribution

4. Domain-Specific

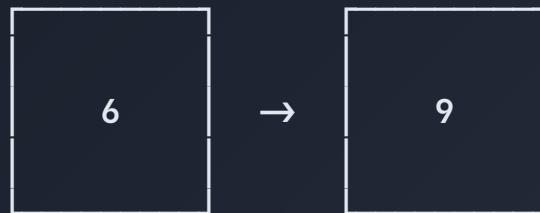# When NOT to Augment

Be careful with:

1. **Medical imaging**: Artifacts can mislead diagnosis

2. **OCR/Text**: Rotation can make text unreadable

3. **Fine-grained classification**: Too much blur loses details

4. **Small objects**: Heavy cropping loses object

5. **Asymmetric objects**: Flips change meaning (e.g., left/right lung)

Rule: Only augment if transformation preserves label

The "6 vs 9" Problem

Classic augmentation mistake:

```
Original:     Flipped Vertically:



   ┌─────┐        ┌─────┐
   │     │        │     │
   │  6  │   →    │  9  │
   │     │        │     │
   └─────┘        └─────┘


Label: "6"      Label: Still "6"???

NO! The label changed! This is WRONG!
```

Always ask: Does this transformation preserve the label?

Good vs Bad Augmentation Examples

GOOD AUGMENTATION:

```
Task: Classify movie genres from posters


Flip horizontal: Action movie is still action
Brightness change: Genre doesn't change
Small rotation: Poster still recognizable
```

BAD AUGMENTATION:

```
Task: Read text from movie posters


Flip horizontal: "STAR WARS" becomes "SRAW RATS"
Heavy rotation: Text becomes unreadable
Too much blur: Can't see letters
```

Text Augmentation: Overview

Challenges:

- Discrete tokens (can't interpolate like pixels)
- Semantic meaning matters
- Grammar and syntax constraints

Approaches:

1. **Rule-based**: Synonym replacement, random operations
2. **Back-translation**: Translate to another language and back
3. **Paraphrasing**: LLMs generate paraphrases
4. **Contextual**: BERT-based word replacement

Text Augmentation: Movie Review Example

```python
# Original review from our Netflix dataset
original = "This movie was absolutely fantastic! Great acting."
label = "POSITIVE"

# Augmented versions (all still POSITIVE)
augmented = [
    "This film was absolutely fantastic! Great acting.",    # Synonym
    "This movie was really fantastic! Great acting.",       # Synonym
    "This movie was absolutely amazing! Great acting.",     # Synonym
    "This movie was fantastic! Excellent acting.",          # Synonym
    "Ce film etait fantastique!" → "This film was great!"   # Back-translation
]

# Now we have 6 training examples from 1!
# All preserve the POSITIVE label
```

Text augmentation must preserve meaning AND sentiment!

The Paraphrase Intuition

Humans express the same idea in many ways:

```
"The movie was great!"
"I really enjoyed this film!"
"Fantastic movie, would recommend!"
"Loved every minute of it!"
"A truly wonderful cinematic experience!"

All mean: POSITIVE sentiment
Model should recognize ALL of these patterns!
```

Text augmentation teaches the model that different words can mean the same thing.

Easy Data Augmentation (EDA) - Simple but effective

4 Operations:

1. **Synonym Replacement**: Replace words with synonyms

```
"The movie was great" → "The film was excellent"
```

2. **Random Insertion**: Insert random synonyms

```
"I love this" → "I really love this"
```

3. **Random Swap**: Swap word positions

```
"She likes pizza" → "She pizza likes"
```

4. **Random Deletion**: Delete words randomly

Text Augmentation with nlpaug

nlpaug: Comprehensive text augmentation library

```python
import nlpaug.augmenter.word as naw
import nlpaug.augmenter.sentence as nas

# Synonym replacement using WordNet
aug_syn = naw.SynonymAug(aug_src='wordnet')
text = "The quick brown fox jumps over the lazy dog"
augmented = aug_syn.augment(text)
print(augmented)
# Output: "The fast brown fox jump over the lazy dog"

# Contextual word embeddings (BERT)
aug_bert = naw.ContextualWordEmbsAug(
    model_path='bert-base-uncased',
    action="substitute"
)
augmented = aug_bert.augment(text)

# Back-translation
aug_back = naw.BackTranslationAug(
    from_model_name='facebook/wmt19-en-de',
    to_model_name='facebook/wmt19-de-en'
)
```

Text Augmentation: Back-Translation

Idea: Translate to another language and back

```python
from transformers import pipeline

# English → German → English
en_de = pipeline("translation", model="Helsinki-NLP/opus-mt-en-de")
de_en = pipeline("translation", model="Helsinki-NLP/opus-mt-de-en")

text = "I love machine learning"
german = en_de(text)[0]['translation_text']
back = de_en(german)[0]['translation_text']

print(f"Original: {text}")
print(f"German: {german}")
print(f"Back: {back}")
# Output: "I love machine learning" → "Ich liebe maschinelles Lernen" → "I love machine learning"
```

Pros: Maintains meaning, natural variations

Cons: Expensive (requires translation models)

Text Augmentation: Paraphrasing with LLMs

Use LLMs to generate paraphrases

```python
from google import genai
import os

client = genai.Client(api_key=os.environ['GEMINI_API_KEY '])

def paraphrase(text, n=3):
    prompt = f"""
    Generate {n} paraphrases of the following text.
    Keep the same meaning but use different words.
    Return one paraphrase per line.

    Text: {text}
    """

    response = client.models.generate_content(
        model="models/gemini-2.0-flash-exp",
        contents=prompt
    )

    paraphrases = response.text.strip().split('\n')
    return paraphrases

text = "The model achieved 95% accuracy"
paraphrases = paraphrase(text, n=3)
for p in paraphrases:
```

Text Augmentation Best Practices

1. Preserve Label

- Sentiment: Don't change positive to negative

- NER: Keep entity boundaries

- Classification: Maintain class meaning

2. Maintain Coherence

- Avoid random operations that break grammar

- Check that output is readable

3. Domain-Specific

- Legal text: Minimal changes (meaning critical)

- Social media: More aggressive OK (informal)

- Code: Very careful (syntax matters)

4. Quality Control

Audio Augmentation: Overview

Audio = Waveform + Spectrogram

Time Domain Augmentations:

- Time stretching

- Pitch shifting

- Adding noise

- Volume changes

- Time shifting

Frequency Domain Augmentations:

- SpecAugment

- Frequency masking

- Time masking

Audio Augmentation with audiomentations

```python
from audiomentations import Compose, AddGaussianNoise, TimeStretch, PitchShift

augment = Compose([
    AddGaussianNoise(min_amplitude=0.001, max_amplitude=0.015, p=0.5),
    TimeStretch(min_rate=0.8, max_rate=1.25, p=0.5),
    PitchShift(min_semitones=-4, max_semitones=4, p=0.5),
])

import librosa

# Load audio
audio, sr = librosa.load('audio.wav', sr=16000)

# Augment
augmented_audio = augment(samples=audio, sample_rate=sr)

# Save
import soundfile as sf
sf.write('augmented_audio.wav', augmented_audio, sr)
```

SpecAugment for Speech Recognition

SpecAugment: Augment spectrograms directly

Operations:

1. **Time Masking**: Mask consecutive time steps

2. **Frequency Masking**: Mask frequency channels

3. **Time Warping**: Warp time axis

```python
import torch
from torchaudio.transforms import FrequencyMasking, TimeMasking

# Convert to spectrogram
spectrogram = torchaudio.transforms.MelSpectrogram()(audio)

# Apply augmentations
freq_mask = FrequencyMasking(freq_mask_param=30)
time_mask = TimeMasking(time_mask_param=100)
```

Used by: Google's speech recognition, Wav2Vec 2.0

Audio Augmentation: Common Techniques

1. Background Noise

```python
from audiomentations import AddBackgroundNoise

augment = AddBackgroundNoise(
    sounds_path="/path/to/noise/files",
    min_snr_db=3,
    max_snr_db=30,
    p=1.0
```

2. Room Impulse Response

```python
from audiomentations import ApplyImpulseResponse

augment = ApplyImpulseResponse(
    ir_path="/path/to/impulse/responses",
    p=0.5
```

3. Compression (MP3 artifacts)
4. Clipping (Simulate distortion)
5. Band-pass filters

Augly: Facebook's Augmentation Library

Unified API for images, audio, video, and text

```python
import augly.image as imaugs
import augly.audio as audaugs
import augly.text as textaugs

# Image
img_augmented = imaugs.augment_image(
    img,
    [
        imaugs.Blur(),
        imaugs.RandomNoise(),
        imaugs.Rotate(degrees=15),
    ]
)

# Audio
audio_augmented = audaugs.apply_lambda(
    audio,
    aug_function=audaugs.add_background_noise,
    snr_level_db=10
)

# Text
text_augmented = textaugs.simulate_typos(
    text,
    aug_char_p=0.05,
    aug_word_p=0.05
```

Augly Features

Cross-Modal Augmentations:

Images:

- Blur, brightness, contrast, noise
- Overlay emoji, text, shapes
- Meme generation
- Pixel distortions

Audio:

- Background noise, reverb, pitch shift
- Clipping, speed, volume
- Time stretch

Text:

- Typos, contractions, slang

Designing an Augmentation Pipeline

Step 1: Understand Your Task

- Classification: Aggressive augmentation OK

- Detection: Preserve bounding boxes

- Segmentation: Transform masks too

Step 2: Start Simple

```
# Baseline: No augmentation
```

Step 3: Gradually Increase

```
transform = A.Compose([
```

Step 4: Measure Impact

- Track validation accuracy

- Compare with/without each augmentation

Augmentation Hyperparameters

Key parameters to tune:

1. **Probability (p)**: How often to apply

   ○ Start: p=0.5

   ○ Increase if underfitting

   ○ Decrease if validation worse

2. **Magnitude**: Strength of transformation

   ○ Rotation: ±10° → ±30°

   ○ Brightness: ±10% → ±30%

3. **Combination**: How many augmentations together

   ○ Start: 1-2 at a time

   ○ Advanced: 3-5 at a time

   Strategy: Grid search or random search

AutoAugment & RandAugment

AutoAugment: Learn augmentation policy with RL

Problem: Manual tuning is tedious

Solution: Use RL to find best augmentation sequence

RandAugment: Simplified version

```python
from torchvision.transforms import RandAugment

transform = RandAugment(
    num_ops=2,          # Number of augmentations to apply
    magnitude=9         # Strength (0-30)
)


augmented = transform(image)
```

Policies learned on ImageNet work well on other datasets!

Test-Time Augmentation (TTA)

Idea: Augment at inference time and average predictions

```python
import albumentations as A

def tta_predict(model, image, n_augments=10):
    transform = A.Compose([
        A.HorizontalFlip(p=0.5),
        A.Rotate(limit=15, p=0.5),
    ])

    predictions = []

    # Original prediction
    predictions.append(model.predict(image))

    # Augmented predictions
    for _ in range(n_augments - 1):
        aug_image = transform(image=image)['image']
        pred = model.predict(aug_image)
        predictions.append(pred)
```

Result: Often 1-2% accuracy improvement

Cost: 10× slower inference

Measuring Augmentation Effectiveness

Experiment Design:

1. **Baseline**: Train without augmentation

2. **With Aug**: Train with augmentation

3. **Compare**:

   ○ Training loss curves

   ○ Validation accuracy

   ○ Test accuracy

   ○ Overfitting gap

```python
# No augmentation
model1 = train_model(train_data, augment=False)
acc_no_aug = model1.evaluate(test_data)

# With augmentation
model2 = train_model(train_data, augment=True)
acc_with_aug = model2.evaluate(test_data)
```

Data Augmentation + Active Learning

Combine both techniques:

1. **Active Learning**: Select most informative samples

2. **Data Augmentation**: Generate variations of selected samples

```python
# Active learning loop with augmentation
for iteration in range(n_iterations):
    # Query uncertain samples
    query_idx = uncertainty_sampling(model, X_pool, n_samples=10)

    # Augment queried samples
    X_aug, y_aug = augment_samples(X_pool[query_idx], y_pool[query_idx])

    # Add original + augmented to training set
    X_train = np.vstack([X_train, X_pool[query_idx], X_aug])
    y_train = np.hstack([y_train, y_pool[query_idx], y_aug])

    # Retrain
```

Result: Best of both worlds!

Domain-Specific Augmentation

Medical Imaging:

- Mild rotations, flips (check anatomy)

- Brightness/contrast (simulate different machines)

- Elastic deformations

- Avoid: Heavy blurs, unrealistic colors

Satellite Imagery:

- Any rotation (no canonical orientation)

- Color shifts (atmospheric conditions)

- Cloud overlays

Document OCR:

- Perspective transforms

- Shadows, creases

Synthetic Data Generation

Beyond augmentation: Generate completely new data

Techniques:

1. **GANs**: Generate realistic images
2. **Style Transfer**: Change image style
3. **3D Rendering**: Render synthetic scenes
4. **Text-to-Image**: Stable Diffusion, DALL-E
5. **Simulation**: Physics engines for robotics

Example: Car detection

- Render 3D car models in various poses
- Add backgrounds
- Train detector

Benefits: Unlimited data, perfect labels

GANs for Data Augmentation

Use trained GAN to generate new samples

```python
from torchvision.models import inception_v3
import torch

# Train GAN on your dataset
# Then generate new samples

generator = load_trained_gan()

# Generate 1000 new images
z = torch.randn(1000, latent_dim)
```

Challenges:

- Training GANs is hard

- May generate unrealistic samples

- Need large dataset to train GAN

Alternative: Use pre-trained GANs (e.g., StyleGAN)

Augmentation for Object Detection

Challenge: Must transform bounding boxes too

```python
import albumentations as A

transform = A.Compose([
    A.HorizontalFlip(p=0.5),
    A.Rotate(limit=10, p=0.5),
    A.RandomBrightnessContrast(p=0.3),
], bbox_params=A.BboxParams(format='pascal_voc', label_fields=['labels']))

# Apply transformation
augmented = transform(
    image=image,
    bboxes=[[23, 45, 120, 150], [50, 80, 200, 250]],
    labels=[0, 1]
)

aug_image = augmented['image']
aug_bboxes = augmented['bboxes']
aug_labels = augmented['labels']
```

Albumentations handles bbox transformations automatically!

Augmentation for Semantic Segmentation

Challenge: Transform masks along with images

```python
transform = A.Compose([
    A.HorizontalFlip(p=0.5),
    A.Rotate(limit=30, p=0.5),
    A.ElasticTransform(p=0.3),
    A.GridDistortion(p=0.3),
])

# Apply to both image and mask
augmented = transform(image=image, mask=mask)

aug_image = augmented['image']
aug_mask = augmented['mask']

# Mask is transformed identically to image
assert aug_image.shape[:2] == aug_mask.shape
```

Common Mistakes

1. Augmenting Test Data

- Only augment training data!
- Test on original distribution

2. Too Strong Augmentation

- Model learns wrong patterns
- Check augmented samples visually

3. Not Preserving Labels

- Digit '6' flipped → '9' (different label!)
- Medical: Left vs right matters

4. Inconsistent Preprocessing

- Normalize before or after augmentation?

Tools & Libraries Summary

Images:

- **Albumentations**: Fast, flexible, comprehensive
- **imgaug**: Similar to Albumentations
- **torchvision.transforms**: PyTorch native
- **Augly**: Facebook's unified library

Text:

- **nlpaug**: Comprehensive text augmentation
- **TextAugment**: EDA implementation
- **Augly.text**: Facebook's text augs

Audio:

- **audiomentations**: Time-domain augmentations
- **torchaudio.transforms**: Frequency-domain

Augmentation in Production

Considerations:

1. **Performance**: Augment on-the-fly vs pre-computed

    ○ On-the-fly: Saves storage, more variety

    ○ Pre-computed: Faster training

2. **Reproducibility**: Set random seeds

```python
random.seed(42)
np.random.seed(42)
torch.manual_seed(42)
```

3. **Validation**: Don't augment val/test sets

4. **Monitoring**: Track which augmentations used

5. **A/B Testing**: Compare models with different augmentations

Research Directions

Current Trends:

1. **Learned Augmentation**: AutoML for augmentation policies

2. **Adversarial Augmentation**: Generate hard examples

3. **Curriculum Augmentation**: Start easy, increase difficulty

4. **Cross-Modal Augmentation**: Transfer between modalities

5. **Foundation Model Augmentation**: Use DALL-E, ChatGPT

Open Problems:

- Optimal augmentation for small datasets

- Task-specific augmentation design

- Augmentation for few-shot learning

- Augmentation quality metrics

Case Study: Image Classification

Dataset: CIFAR-10 (10 classes, 50k train images)

Baseline (No Augmentation):

- Train accuracy: 99%
- Test accuracy: 70%
- Clear overfitting!

With Standard Augmentation:

```python
transform = A.Compose([
    A.HorizontalFlip(p=0.5)
```

- Train accuracy: 85%
- Test accuracy: 82%
- Better generalization!

Improvement: +12% test accuracy

What We've Learned

Core Concepts:

- Data augmentation creates training data variations

- Preserves labels while increasing diversity

- Reduces overfitting and improves generalization

Techniques:

- Image: Geometric + color transforms

- Text: Synonym replacement, back-translation, paraphrasing

- Audio: Time stretching, pitch shifting, noise

Libraries:

- Albumentations (images)

- nlpaug (text)

- audiomentations (audio)

Practical Recommendations

Getting Started:

1. Use Albumentations for images

2. Start with flip + rotate + brightness

3. Measure baseline vs augmented

4. Gradually add more augmentations

Hyperparameter Tuning:

- Probability: 0.3-0.7

- Magnitude: Start low, increase if underfitting

- Number of augs: 2-4 simultaneously

Production:

- Augment on-the-fly during training

Resources

Papers:

- "AutoAugment: Learning Augmentation Policies from Data" (2019)

- "RandAugment: Practical automated data augmentation" (2020)

- "SpecAugment: A Simple Data Augmentation Method for ASR" (2019)

- "mixup: Beyond Empirical Risk Minimization" (2018)

Libraries:

- Albumentations: https://albumentations.ai/

- nlpaug: https://github.com/makcedward/nlpaug

- audiomentations: https://github.com/iver56/audiomentations

- Augly: https://github.com/facebookresearch/AugLy

Tutorials:

- Albumentations documentation

Mathematical Foundations: Invariance and Equivariance

Invariance: Output doesn't change under transformation

$$f(T(x)) = f(x)$$

Example: Image classifier should be invariant to rotation

- $f(\text{rotate}(\text{cat})) = \text{"cat"}$

Equivariance: Output transforms consistently with input

$$f(T(x)) = T'(f(x))$$

Example: Segmentation should be equivariant to rotation

- $\text{segment}(\text{rotate}(\text{image})) = \text{rotate}(\text{segment}(\text{image}))$

Data augmentation teaches invariance:

```python
# Training with augmented data
loss = cross_entropy(model(rotate(x)), y)   # Same label y
# Model learns rotation invariance
```

Manifold Hypothesis and Augmentation

Manifold Hypothesis: High-dimensional data lies on low-dimensional manifold

Augmentation explores the manifold:

- Original data: Sparse samples on manifold
- Augmented data: Fill gaps along manifold

Interpolation on manifold:

$$x_{aug} = x + \epsilon \cdot \nabla_\theta T(x)$$

where T is transformation, ϵ is small

Theoretical benefit:

- Smoother decision boundaries
- Better generalization
- Reduced sample complexity

Mixup: Theory and Implementation

Mixup: Linear interpolation of examples and labels

Formula:

$$\tilde{x} = \lambda x_i + (1 - \lambda)x_j$$
$$\tilde{y} = \lambda y_i + (1 - \lambda)y_j$$

where $\lambda \sim \text{Beta}(\alpha, \alpha)$

Theoretical motivation:

- Vicinal Risk Minimization (VRM)
- Encourages linear behavior between training examples
- Regularizes network to output convex combinations

Implementation:

```python
def mixup_data(x, y, alpha=1.0):
    """Mixup data and labels."""
    if alpha > 0:
```

Mixup Variants: CutMix and MoEx

CutMix: Replace patches instead of blending

Advantages over Mixup:

- Preserves localization ability
- More efficient for CNNs (no blend artifacts)

```python
def cutmix(x, y, alpha=1.0):
    """CutMix augmentation."""
    lam = np.random.beta(alpha, alpha)
    batch_size, _, H, W = x.shape
    index = torch.randperm(batch_size)

    # Random box
    cut_rat = np.sqrt(1. - lam)
    cut_w = int(W * cut_rat)
    cut_h = int(H * cut_rat)

    cx = np.random.randint(W)
    cy = np.random.randint(H)
```

MoEx (Momentum Exchange): Exponential moving average blending

Diffusion Models for Data Augmentation

Modern approach: Use diffusion models to generate variations

Workflow:

1. Add small noise to image
2. Denoise with pretrained diffusion model
3. Use denoised version as augmentation

```python
from diffusers import StableDiffusionImg2ImgPipeline
```

Benefits:

- Semantically meaningful variations
- Controllable via prompts
- High quality

Challenges: Expensive, requires GPU

Contrastive Learning Augmentation Strategies

SimCLR: Self-supervised learning via contrastive loss

Key idea: Different augmentations of same image should have similar representations

Augmentation composition:

```python
import torchvision.transforms as transforms

# SimCLR augmentation pipeline
simclr_transform = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.RandomApply([
        transforms.ColorJitter(0.8, 0.8, 0.8, 0.2)
```

Findings:

- Crop + color jitter most important

- Composition matters more than individual augmentations

- Stronger augmentation → better representations

59

MoCo (Momentum Contrast) Augmentation

MoCo v2 augmentation:

```python
moco_transform = transforms.Compose([
    transforms.RandomResizedCrop(224, scale=(0.2, 1.0)),
    transforms.RandomApply([
        transforms.ColorJitter(0.4, 0.4, 0.4, 0.1)
    ], p=0.8),
    transforms.RandomGrayscale(p=0.2),
    transforms.RandomApply([transforms.GaussianBlur(kernel_size=23)], p=0.5),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

Queue-based approach:

- Maintain queue of negatives

- Momentum encoder for consistency

Result: State-of-art self-supervised learning

Invariant Risk Minimization (IRM)

Goal: Learn invariant features across environments

Formulation:

$$\min_{\Phi} \sum_{e \in \mathcal{E}} R^e(\Phi) + \lambda \|\nabla_{w|w=1.0} R^e(w \cdot \Phi)\|^2$$

where:

- $\mathcal{E}$: Set of environments (different augmentations)
- R^e: Risk in environment e
- Φ: Feature extractor

Augmentation as environments:

```python
def irm_loss(model, x, y, augmentations):
    """IRM loss across augmentation environments."""
    total_loss = 0
```

Benefit: Robust to distribution shift

Consistency Regularization: UDA and FixMatch

Unsupervised Data Augmentation (UDA):

Idea: Model predictions should be consistent under augmentation

$$L_{consistency} = \mathbb{E}_{x,aug}[KL(p(y|x)\|p(y|aug(x)))]$$

```python
def uda_loss(model, x_unlabeled, strong_aug, weak_aug):
    """UDA consistency loss."""
    # Weak augmentation prediction (pseudo-label)
    with torch.no_grad():
        weak_pred = model(weak_aug(x_unlabeled))
        pseudo_label = torch.softmax(weak_pred, dim=1)
```

FixMatch: UDA + pseudo-labeling with confidence threshold

```python
def fixmatch_loss(model, x_unlabeled, threshold=0.95):
    """FixMatch semi-supervised loss."""
    # Weak augmentation
    weak_pred = model(weak_aug(x_unlabeled))
    max_probs, pseudo_labels = torch.max(torch.softmax(weak_pred, dim=1), dim=1)
```

Learnable Augmentation Policies

Neural Augmentation: Learn transformation parameters

Approach:

```python
class LearnableAugmentation(nn.Module):
    def __init__(self):
        super().__init__()
        # Learnable parameters for augmentation
        self.rotation_range = nn.Parameter(torch.tensor(15.0))
        self.brightness_factor = nn.Parameter(torch.tensor(0.2))

    def forward(self, x):
        # Apply augmentation with learned parameters
        angle = torch.rand(1) * self.rotation_range
        brightness = 1 + torch.rand(1) * self.brightness_factor

        x_aug = rotate(x, angle)
        x_aug = adjust_brightness(x_aug, brightness)

        return x_aug

# Training: Backprop through augmentation
learnable_aug = LearnableAugmentation()
optimizer = torch.optim.Adam(learnable_aug.parameters())
```

Benefit: Automatically tune augmentation strength

Adversarial Training as Augmentation

Adversarial examples: Inputs with small perturbations that fool model

PGD (Projected Gradient Descent):

$$x_{adv} = x + \epsilon \cdot \text{sign}(\nabla_x L(f(x), y))$$

Adversarial training:

```python
def pgd_attack(model, x, y, epsilon=0.3, alpha=0.01, num_iter=10):
    """Generate adversarial example."""
    x_adv = x.clone().detach()

    for _ in range(num_iter):
        x_adv.requires_grad = True

        # Compute loss
        pred = model(x_adv)
        loss = F.cross_entropy(pred, y)

        # Gradient ascent
        loss.backward()
        grad = x_adv.grad

        # Update adversarial example
        x_adv = x_adv + alpha * grad.sign()

        # Project back to epsilon ball
        perturbation = torch.clamp(x_adv - x, -epsilon, epsilon)
        x_adv = torch.clamp(x + perturbation, 0, 1).detach()

    return x_adv

# Training with adversarial examples
for x, y in dataloader:
```

Meta-Learning for Augmentation

Goal: Learn which augmentations help for specific tasks

Meta-augmentation:

```python
class MetaAugmentation:
    def __init__(self, augmentations):
        self.augmentations = augmentations
        # Learnable weights for each augmentation
        self.weights = nn.Parameter(torch.ones(len(augmentations)))

    def sample_augmentation(self):
        """Sample augmentation based on learned weights."""
        probs = F.softmax(self.weights, dim=0)
        idx = torch.multinomial(probs, 1).item()
        return self.augmentations[idx]

    def meta_train(self, meta_train_tasks, meta_val_tasks):
        """Meta-training loop."""
        for epoch in range(num_epochs):
            for task in meta_train_tasks:
                # Sample augmentation
                aug = self.sample_augmentation()

                # Train on augmented data
                x_aug, y_aug = aug(task.x_train), task.y_train
                model.train_step(x_aug, y_aug)
```

Benefit: Task-specific augmentation policies

Augmentation Budget and Efficiency

Computational cost:

Augmentation	Cost (ms/image)	Speedup Strategy
Horizontal flip	0.1	Already fast
Rotation	2.0	Use smaller angles
Color jitter	1.5	GPU acceleration

Optimization strategies:

```python
# 1. GPU acceleration
import kornia

transform = kornia.augmentation.AugmentationSequential(
    kornia.augmentation.RandomRotation(30),
    kornia.augmentation.ColorJitter(0.2, 0.2, 0.2, 0.1),
    data_keys=["input"]
)

# Apply on GPU (batched)
x_aug = transform(x_gpu)   # Much faster than CPU

# 2. Caching expensive augmentations
class CachedAugmentation:
    def __init__(self, aug_fn, cache_size=10000):
```

Data Mixing Beyond Mixup

SaliencyMix: Mix based on saliency maps

```python
def saliencymix(x, y, saliency_fn):
    """Mix based on saliency."""
    batch_size = x.size(0)
    index = torch.randperm(batch_size)

    # Get saliency maps
    sal_a = saliency_fn(x)
    sal_b = saliency_fn(x[index])

    # Mix based on saliency
    mask = (sal_a > sal_b).float()
```

PuzzleMix: Mix by solving optimization problem

- Find optimal cut to preserve features

- More sophisticated than random cuts

Co-Mixup: Mix within same class to preserve fine-grained features

Policy Search for Optimal Augmentation

Population Based Augmentation (PBA):

Algorithm:

1. Initialize population of augmentation policies

2. Train models with different policies

3. Select best performers

4. Mutate and combine policies

5. Repeat

```python
class AugmentationPolicy:
    def __init__(self):
        self.ops = random.sample(ALL_OPS, k=2)
        self.probs = np.random.uniform(0, 1, size=2)
        self.magnitudes = np.random.uniform(0, 1, size=2)

    def mutate(self):
        """Mutate policy."""
        idx = random.randint(0, 4)
        if random.random() < 0.5:
            self.probs[idx] += np.random.normal(0, 0.1)
        else:
            self.magnitudes[idx] += np.random.normal(0, 0.1)

    def crossover(self, other):
        """Combine two policies."""
        child = AugmentationPolicy()
        for i in range(5):
            if random.random() < 0.5:
                child.ops[i] = self.ops[i]
```

Augmentation for Long-Tail Distribution

Problem: Rare classes benefit more from augmentation

Class-balanced augmentation:

```python
class ClassBalancedAugmentation:
    def __init__(self, class_counts):
        # Compute augmentation probability per class
        # More augmentation for rare classes
        total = sum(class_counts)
        self.aug_probs = {
            cls: 1.0 - (count / total)
            for cls, count in enumerate(class_counts)
        }

    def __call__(self, x, y):
        """Apply augmentation based on class."""
        aug_prob = self.aug_probs[y]

        if random.random() < aug_prob:
            # Strong augmentation for rare classes
            x = strong_augment(x)
        else:
            # Weak augmentation for common classes
```

Remix: Oversample tail classes with mixup

BBN: Bilateral-branch network with different augmentation per branch

Temporal Augmentation for Videos

Video-specific challenges:

- Temporal consistency

- Motion patterns

- Longer sequences

Temporal augmentation:

```python
def temporal_augment(video, fps=30):
    """Augment video data."""
    # 1. Temporal crop
    start = random.randint(0, len(video) - 64)
    video = video[start:start+64]
```

Spatial + Temporal:

- Apply same spatial aug to all frames (consistency)

- Or different augs per frame (diversity)

3D Augmentation for Point Clouds

Point cloud augmentation:

```python
def pointcloud_augment(points):
    """Augment 3D point cloud."""
    # 1. Random rotation
    angle = np.random.uniform(0, 2*np.pi)
    rotation_matrix = np.array([
        [np.cos(angle), -np.sin(angle), 0],
        [np.sin(angle), np.cos(angle), 0],
        [0, 0, 1]
    ])
    points = points @ rotation_matrix.T

    # 2. Random scaling
    scale = np.random.uniform(0.8, 1.2)
    points = points * scale

    # 3. Random jitter
    noise = np.random.normal(0, 0.02, size=points.shape)
    points = points + noise

    # 4. Random point dropout
```

PointAugment: Learnable augmentation for point clouds

PointMixup: Mixup in 3D space

Graph Augmentation for GNNs

Graph-specific augmentation:

```python
def graph_augment(graph):
    """Augment graph structure."""
    # 1. Edge dropping
    edge_mask = torch.rand(graph.num_edges) > 0.1
    graph.edge_index = graph.edge_index[:, edge_mask]

    # 2. Node dropping
    node_mask = torch.rand(graph.num_nodes) > 0.1
    graph = graph.subgraph(node_mask)

    # 3. Feature masking
    feat_mask = torch.rand(graph.x.size(1)) > 0.2
    graph.x[:, ~feat_mask] = 0

    # 4. Edge perturbation (add random edges)
    n_new_edges = int(0.1 * graph.num_edges)
    src = torch.randint(0, graph.num_nodes, (n_new_edges,))
    dst = torch.randint(0, graph.num_nodes, (n_new_edges,))
```

GraphCL: Contrastive learning for graphs with augmentation

M-Mix: Mixup for molecular graphs

Augmentation Evaluation Metrics

How to measure augmentation quality?

1. Downstream Performance:

```python
def evaluate_augmentation(aug_fn, model, data):
    """Evaluate by downstream task performance."""
    # Train with augmentation
    model_aug = train_model(data, augmentation=aug_fn)
```

2. Diversity Score:

```python
def diversity_score(original, augmented):
    """Measure diversity of augmented samples."""
    # Compute pairwise distances
```

3. Invariance Test:

```python
def invariance_score(model, x, augmentations):
    """Measure how invariant model is to augmentations."""
    original_pred = model(x)
```

Curriculum Augmentation

Idea: Start with weak augmentation, gradually increase strength

Progressive augmentation:

```python
class CurriculumAugmentation:
    def __init__(self, max_epochs):
        self.max_epochs = max_epochs
        self.current_epoch = 0

    def get_augmentation(self):
        """Return augmentation based on training progress."""
        # Linearly increase augmentation strength
        progress = self.current_epoch / self.max_epochs

        if progress < 0.3:
            # Early: weak augmentation
            return A.Compose([
                A.HorizontalFlip(p=0.5),
            ])
        elif progress < 0.7:
            # Mid: medium augmentation
            return A.Compose([
                A.HorizontalFlip(p=0.5),
                A.Rotate(limit=15, p=0.5),
                A.RandomBrightnessContrast(p=0.3),
            ])
        else:
            # Late: strong augmentation
            return A.Compose([
                A.HorizontalFlip(p=0.5),
                A.Rotate(limit=30, p=0.5),
                A.RandomBrightnessContrast(p=0.5),
                A.GaussNoise(p=0.3),
                A.Cutout(num_holes=8, max_h_size=16, max_w_size=16, p=0.5),
            ])

    def update_epoch(self, epoch):
        self.current_epoch = epoch

# Training loop
curr_aug = CurriculumAugmentation(max_epochs=100)

for epoch in range(100):
    augmentation = curr_aug.get_augmentation()
```

Multi-Modal Augmentation

Cross-modal augmentation: Augment multiple modalities consistently

Example: Image + Text

```python
def multimodal_augment(image, caption):
    """Augment image and caption together."""
    # Image augmentation
    if random.random() < 0.5:
        image = horizontal_flip(image)
        # Update caption if needed
        # "person on left" → "person on right"
        caption = flip_spatial_words(caption)
```

Audio + Text (speech recognition):

```python
def audio_text_augment(audio, transcript):
    """Augment audio and transcript together."""
    # Speed perturbation
    speed = random.uniform(0.9, 1.1)
    audio = change_speed(audio, speed)
    # Transcript unchanged (same words)
```

Foundation Model-Based Augmentation

Stable Diffusion for augmentation:

```python
from diffusers import StableDiffusionPipeline

pipe = StableDiffusionPipeline.from_pretrained("stabilityai/stable-diffusion-2-1")

def foundation_augment(original_image, class_name, n_augments=5):
```

Benefits:

- Semantically meaningful variations
- Can generate rare classes
- High visual quality

Challenges:

- Expensive (GPU, time)
- May generate out-of-distribution samples
- Need careful prompt engineering

Augmentation Transferability

Question: Do augmentations learned on one dataset transfer to others?

Empirical findings:

ImageNet → Other Vision Tasks:

- AutoAugment policies from ImageNet work well on CIFAR, SVHN
- **Transferability**: ~80-90% of performance

Natural Images → Medical Images:

- Standard augmentations (rotation, flip) transfer well
- Advanced (CutMix, MixUp) less effective
- **Transferability**: ~60-70%

Practical implications:

```
# Use pre-discovered policies for your domain
```

Advanced Augmentation Summary

Theoretical Foundations:

- Invariance and equivariance

- Manifold hypothesis

- Vicinal risk minimization (Mixup)

- Consistency regularization

Advanced Mixing Strategies:

- Mixup, CutMix, MoEx

- SaliencyMix, PuzzleMix

- Class-balanced mixing

Modern Approaches:

- Diffusion models for augmentation

- Contrastive learning (SimCLR, MoCo)

Questions?

Lab: Implement and compare augmentation strategies

Measure impact on model performance